

Milestone 4 - System Prototype

Team GitHub notebook URL: https://github.com/Kongbai815/job-matching-nlp/blob/main/MSBA_Job_Matching_Milestone4_System_Prototype.ipynb

1. End-to-End Architecture

The system separates search requests from posting classification, retrieves comparable real postings, scores postings with a trained word+character TF-IDF SVM, applies policy gates, and generates an evidence-bound advisor response.

Stage	Implementation	Output
Route	search vs. posting input mode	Correct task semantics
Retrieve	80k sparse index + constraint boosts + dedupe	Top-6 evidence
Classify	LinearSVC trained on 77,135 deduplicated rows	Four-class role fit
Govern	missing-field + authorization gates	Review / Hold
Generate	deterministic evidence template	Citations and caveat

2. Executed Setup

```
from msba_job_matcher.core import JobMatchingSystem
system = JobMatchingSystem(DATA, model_path=MODEL)

Project rows: 100000
Retrieval rows: 80000
Classifier: word+character TF-IDF LinearSVC (77,135 deduplicated training rows)
```

3. Workflow Inputs

Case	Mode	Model label	System label	Review
entry level analytics search	search	N/A (search request)	N/A (search request)	True
senior engineering posting	posting	low_fit	low_fit	True
authorization restriction posting	posting	high_fit	high_fit	True
thin metadata posting	posting	medium_fit	unclear	True

4. Retrieved Evidence and Policy Behavior

Case	Top evidence	Country	Authorization / policy
entry level analytics search	Data Analyst	United States	not available in public source
senior engineering posting	Data Engineer	United States	not available in public source
authorization restriction posting	Jr. Data Analyst (No OPT/CPT)	United States	explicit source text: no OPT/CPT
thin metadata posting	Summer Data Analytics Internship	United States	Missing: company, location, country, schedule

5. Evaluation Against Baseline

System	Rows	Accuracy	Macro F1
M2 same-subset baseline	4,000	-	0.791
M3 true PEFT LoRA	4,000	0.875	0.873
M4 trained classifier	4,000	0.984	0.984

Same-subset macro-F1 improvement: **+0.192**. All four final per-label F1 scores are at least 0.970.

6. Leakage-Controlled Checks

Validation slice	Rows	Macro F1
Fixed set, novel text	3,751	0.983
Full validation, novel text	18,672	0.979
Full validation, unseen companies	5,570	0.980
Newest 20% by posting date	20,000	0.976

7. Grounding, Hallucination, and Governance

All four cases returned retrieved evidence. The authorization case demonstrates the critical conflict: the classifier predicts high role fit, but explicit No OPT/CPT text forces a Hold and advisor review. The thin-metadata case demonstrates a second policy override: missing company, location, country, and schedule fields route the result to unclear. The saved casebook contains zero unsupported authorization claims.

8. What the Model Comparison Means

A controlled learning curve rises from 0.898 macro F1 at 2,000 rows to 0.984 at 80,000 rows. The final linear classifier also exceeds the 2,000-row LoRA result while producing a 3.8 MB artifact and training in under one minute. The time-based test confirms that the result is not limited to a random split.

9. Completed Validation Assets

The project scanned all 785,741 source rows and found 230 unique title-level authorization candidates. It includes a 430-row candidate/control benchmark and a 1,000-row advisor queue: 800 low-margin role-fit cases plus 200 authorization candidates. Human judgment fields are intentionally blank so no model output is presented as advisor ground truth.

10. Reading Connection

HOLLM Chapter 8 motivates retrieval as external grounding. Tunstall Chapter 7 and Jurafsky & Martin Chapter 11 emphasize supported QA and extraction. The prototype follows that logic by keeping search, posting classification, retrieved evidence, missing-data policy, and authorization evidence as separate inspectable stages.

Milestone 4 Technical Summary

Architecture and End-to-End Behavior

The system now implements two real workflows. Search mode retrieves and ranks opportunities without pretending the search sentence is a posting. Posting mode classifies a job description, retrieves comparable evidence, and applies deterministic missing-data and authorization gates.

Layer	Method	Business control
Mode	search vs. posting	Prevents task confusion
Retrieval	80k sparse index + boosts + dedupe	Inspectable candidates
Decision	word+character TF-IDF LinearSVC	0.984 macro F1
Policy	missing fields + authorization evidence	Unclear / Hold
Response	evidence-bound template	No unsupported facts

Evaluation vs. Baseline

System	Rows	Accuracy	Macro F1
M2 same-subset baseline	4,000	-	0.791
M3 true PEFT LoRA	4,000	0.875	0.873
M4 trained decision layer	4,000	0.984	0.984

The same-subset macro-F1 gain is +0.192. The model was trained on 77,135 rows after exact-text deduplication, not only the 2,000 rows used by M3 LoRA.

Robustness Audit

The original random split contains 1,328 exact validation texts that also appear in training. After controlling for this, macro F1 remains 0.979 on 18,672 novel-text rows, 0.980 on 5,570 rows from unseen companies, and 0.976 on the newest 20,000 postings.

Grounded Workflow Evidence

Case	System result	Operational response
entry level analytics search	N/A (search)	Advisor review before forwarding
senior engineering posting	low_fit	Advisor review before forwarding
authorization restriction posting	high_fit	Hold - explicit authorization restriction requires advisor verification
thin metadata posting	unclear	Advisor review before forwarding

Risk and Grounding Analysis

Risk 1 is authorization misinformation. Role fit and authorization are separate: explicit negative language forces a Hold, and missing evidence never becomes eligibility. Risk 2 is weak-label overconfidence. The high score measures agreement with transparent rules, so the system preserves human review and reports novel-text and unseen-company slices rather than relying on a random split alone.

Completed Validation Package

The repository now includes a full 785,741-row source scan, 230 unique title-level authorization candidates, a 430-row candidate/control benchmark, and a 1,000-row advisor queue. The queue balances 800 low-margin role-fit cases with 200 authorization candidates; its human fields remain blank for genuine advisor review.

Deployment Boundary

Use the prototype for first-pass prioritization. Search results, low-margin classifications, missing fields, and any authorization-sensitive case remain advisor-confirmed. No automated forwarding, rejection, or immigration decision is permitted.