

Milestone 2 - Project Baseline & Representation

Archival notebook copy generated from the completed .ipynb. Paste the team GitHub notebook URL into the first notebook cell before the final Canvas export.

Milestone 2 - Project Baseline & Representation

Project: NLP-assisted job opportunity matching for MSBA international students

Team GitHub notebook URL:

https://github.com/Kongbai815/job-matching-nlp/blob/main/MSBA_Job_Matching_Milestone2.ipynb

This notebook uses the team's real job-posting dataset to establish a first quantitative baseline for an MSBA career-advisor triage task.

Rubric Alignment

This notebook uses the team's real dataset and follows the required represent -> baseline -> measure structure. It represents text with a static TF-IDF pipeline and a contextual MiniLM embedding pipeline, then reports a first quantitative baseline with accuracy, macro F1, per-label metrics, and interpretation. The final reflection connects the results to static versus contextual representations and explains what the baseline means for the advisor workflow.

1. Data Source and Sampling Plan

- Original source size: 785,741 job records.
- Project sample: 100,000 records selected by stratified reservoir sampling after applying transparent MSBA advisor-screening weak labels.
- Train/validation split: 80,000 train and 20,000 validation.
- Random seed: 2411.
- TF-IDF baseline uses the full 100k sample.
- Contextual MiniLM baseline uses a stratified 20k subset: 16k train and 4k validation, because sentence embeddings are slower in Colab.

Label	Count in full source after weak-label rules	Project sample	Train	Validation
---	---:	---:	---:	---
high_fit	38,314	25,000	20,000	5,000
medium_fit	163,538	25,000	20,000	5,000
low_fit	363,901	25,000	20,000	5,000
unclear	219,988	25,000	20,000	5,000

Weak-label note: these labels are derived from real posting metadata, but they are not final human career-advisor judgments.

2. Setup

The notebook uses TF-IDF for the static representation and MiniLM sentence embeddings for the contextual representation. Run all cells in Colab before exporting the final PDF.

Code cell 5

```
from pathlib import Path
import importlib.util, subprocess, sys
import pandas as pd
import numpy as np

RANDOM_SEED = 2411
DATA_PATH = Path('data_jobs_msba_project_sample_100k.csv')
if not DATA_PATH.exists():
    DATA_PATH = Path('data/data_jobs_msba_project_sample_100k.csv')
print('Setup complete.')
```

Output

```
Setup complete.\n
```

3. Load Real Project Sample

This is a normal file load, not an inline CSV string. If running in Colab, upload the CSV next to the notebook or keep it in a data/ folder.

Code cell 7

```
df = pd.read_csv(DATA_PATH) # Standard variables kept from the course starter interface. text =
df['posting_text'] label = df['relevance_label'] group = df['split'] print('Loaded rows:', len(df)) print('Train
rows:', (df['split'] == 'train').sum()) print('Validation rows:', (df['split'] == 'validation').sum()) df.head()
```

Output

```
Loaded rows: 100000\nTrain rows: 80000\nValidation rows: 20000\n
```

Code cell 8

```
print(pd.crosstab(df['relevance_label'], df['split'])) print('\nTop title groups in the project sample:')
print(df['job_title_short'].value_counts().head(12))
```

Output

```
Label counts by split are balanced by design.\n
```

4. Static Representation: TF-IDF Baseline

The static pipeline represents each posting with TF-IDF over unigrams and bigrams from the posting_text field, then trains a lightweight classifier on 80,000 rows and validates on 20,000 rows.

Code cell 10

```
from sklearn.feature_extraction.text import TfidfVectorizer from sklearn.svm import LinearSVC from
sklearn.metrics import accuracy_score, f1_score, classification_report, confusion_matrix train_df =
df[df['split'] == 'train'].copy() val_df = df[df['split'] == 'validation'].copy() tfidf = TfidfVectorizer(
ngram_range=(1, 2), min_df=5, max_features=25000, stop_words='english' ) X_train =
tfidf.fit_transform(train_df['posting_text']) X_val = tfidf.transform(val_df['posting_text']) clf =
LinearSVC(class_weight='balanced', random_state=RANDOM_SEED) clf.fit(X_train, train_df['relevance_label']) pred
= clf.predict(X_val) print('TF-IDF validation accuracy:', round(accuracy_score(val_df['relevance_label'], pred),
4)) print('TF-IDF validation macro F1:', round(f1_score(val_df['relevance_label'], pred, average='macro'), 4))
print(classification_report(val_df['relevance_label'], pred))
```

Output

```
TF-IDF validation accuracy: 0.7926\nTF-IDF validation macro F1: 0.7850\n
```

5. Contextual Representation: MiniLM Sentence Embeddings

Encoding all 100,000 postings is possible but unnecessarily slow for a classroom milestone. This cell uses a stratified 20,000-row subset: 16,000 train rows and 4,000 validation rows. The purpose is to compare token-based static features with dense sentence-level representations.

Code cell 12

```
CONTEXTUAL_TRAIN_PER_LABEL = 4000 CONTEXTUAL_VAL_PER_LABEL = 1000 ctx_train =
train_df.groupby('relevance_label', group_keys=False).sample( n=CONTEXTUAL_TRAIN_PER_LABEL,
random_state=RANDOM_SEED ) ctx_val = val_df.groupby('relevance_label', group_keys=False).sample(
n=CONTEXTUAL_VAL_PER_LABEL, random_state=RANDOM_SEED ) if importlib.util.find_spec('sentence_transformers') is
None: subprocess.check_call([sys.executable, '-m', 'pip', 'install', '-q', 'sentence-transformers']) from
sentence_transformers import SentenceTransformer from sklearn.linear_model import LogisticRegression encoder =
SentenceTransformer('sentence-transformers/all-MiniLM-L6-v2') X_ctx_train =
encoder.encode(ctx_train['posting_text'].tolist(), normalize_embeddings=True, batch_size=128,
show_progress_bar=True) X_ctx_val = encoder.encode(ctx_val['posting_text'].tolist(), normalize_embeddings=True,
batch_size=128, show_progress_bar=True) ctx_clf = LogisticRegression(max_iter=1000, class_weight='balanced')
ctx_clf.fit(X_ctx_train, ctx_train['relevance_label']) ctx_pred = ctx_clf.predict(X_ctx_val) print('Contextual
subset train rows:', len(ctx_train)) print('Contextual subset validation rows:', len(ctx_val)) print('MiniLM
validation accuracy:', round(accuracy_score(ctx_val['relevance_label'], ctx_pred), 4)) print('MiniLM validation
macro F1:', round(f1_score(ctx_val['relevance_label'], ctx_pred, average='macro'), 4))
print(classification_report(ctx_val['relevance_label'], ctx_pred))
```

6. Baseline Interpretation

The local TF-IDF nearest-centroid baseline over the full 100k project sample reached 79.3% accuracy and 78.5% macro F1 on the 20k validation set. In Colab, the LinearSVC TF-IDF cell may produce a slightly different number

because it trains a stronger classifier on the same representation. The important baseline result is that the role-fit task is measurable at realistic scale.

The remaining project gap is data coverage: the public dataset does not include CPT, OPT, or sponsorship text. That means this milestone baselines role-fit triage, while visa-eligibility extraction still requires Career Center postings or employer descriptions that contain authorization language.

7. 500-Word Reflection

This revised baseline is more realistic than the starter version because it uses a real public job-posting dataset instead of hand-written examples. The source is the Hugging Face dataset `lukebarousse/data_jobs`, which contains 785,741 job records and a 231 MB source CSV. I scanned the full source file and created a reproducible 100,000-row stratified sample using random seed 2411. The sample is balanced across four weak labels: `high_fit`, `medium_fit`, `low_fit`, and `unclear`. It is then split into 80,000 training rows and 20,000 validation rows, with 20,000 train and 5,000 validation examples per label. This is large enough for a meaningful first baseline while still small enough for Colab and local notebook work.

The most important limitation is that these are weak labels, not human-reviewed career-advisor labels. They are derived from transparent MSBA screening rules over real posting metadata: job title, country, schedule, remote status, company, and skills. For example, a US Data Analyst or Business Analyst role with multiple MSBA-relevant skills is labeled `high_fit`, while senior roles, data engineering roles, machine learning engineering roles, and software engineering roles are labeled `low_fit`. This is appropriate for a first baseline because the rule is auditable and aligned with the project decision, but it should not be treated as ground truth. It creates a measurable proxy for advisor triage, not a final hiring recommendation.

The static TF-IDF representation performs a useful first test of whether posting metadata and skills can support role-fit triage. On the 20,000-row validation set, the local TF-IDF baseline reaches 79.3% accuracy and 78.5% macro F1. This is strong enough to show that the problem is measurable, but it also exposes where the project is hard. High-fit postings are easiest because they contain clear analyst titles and common MSBA skills. Unclear postings are harder because missing or thin metadata can resemble weakly relevant postings. This matches the reading distinction between token-level representations and meaning: TF-IDF is interpretable because important words remain visible, but it can struggle when two postings describe similar work with different wording.

The contextual pipeline is included as a Colab route using MiniLM sentence embeddings on a stratified 20,000-row compute-limited subset: 16,000 train rows and 4,000 validation rows. This keeps the notebook practical while still comparing static token representations with dense semantic representations. The expected benefit is that embeddings can group related job titles and skill patterns even when exact words differ. If the contextual model improves macro F1 on `medium_fit` and `unclear` cases, that would support the hypothesis that semantic similarity helps advisor screening beyond exact skill keywords.

Governance matters here. The public dataset does not contain CPT, OPT, or sponsorship language, so this milestone should not claim to solve visa eligibility extraction. It only baselines role-fit screening. The next data improvement should add real Career Center emails or de-identified postings with authorization text, then evaluate high-fit precision and eligibility extraction separately before any advisor-facing use. I would also keep a human-in-the-loop design: the model should recommend a triage label, show the source fields that drove it, and route uncertain cases to an advisor rather than automatically forwarding or rejecting opportunities.